

EntityWorks Component Relationship Map — Structural Reference (v1.1)

Human-Facing Structural Reference

Document Title: EntityWorks Component Relationship Map — Structural Reference

Edition: Human-Facing Structural Reference

Version: v1.1

Status: Published

Publication Date: 27 December 2025

Maintained by: EntityWorks Ltd

Abstract

This document provides a structured description of how the published components of the EntityWorks body of work relate to one another. It sets out the positional relationships between definitions, architectural elements, boundary artefacts, diagnostic constructs, evaluative components, analytical functions, and signalling mechanisms as they appear across existing EntityWorks publications.

The purpose of this reference is to support consistent interpretation of how individual components are situated relative to one another, without restating their internal definitions or duplicating material published elsewhere.

This document is descriptive in nature and does not introduce new components, functions, or mechanisms.

Scope and Status Notice

This document applies only to the components explicitly defined and published by EntityWorks at the time of release. It describes relationships between existing components as documented in their respective publications.

This reference does not define new standards, prescribe implementation approaches, combine component functions, or extend the scope of any individual component. It does not describe system behaviour, technical implementations, or internal mechanisms of AI systems.

EntityWorks Component Relationship Map

Structural Reference (v1.1)

1. Purpose of This Reference

This document provides a **formal structural reference** defining how the core components of the EntityWorks Standard relate to one another. It establishes the canonical relational topology through which the Standard is interpreted, ensuring that each component is understood according to its defined role and position.

The purpose of this reference is not to introduce new mechanisms, requirements, or evaluative criteria. Instead, it exists to **preserve structural clarity** by making explicit the separation of concerns across the Standard. By defining how components connect—and, critically, how they do *not*—this document prevents misinterpretation, functional overlap, and role collapse as the Standard is applied, extended, or discussed.

2. Structural Principle: Layered Representational Topology

The EntityWorks Standard is organised as a **layered representational system**, not as a single framework or unified mechanism. Each layer addresses a distinct aspect of how entities are represented, interpreted, evaluated, and signalled, and each depends conceptually on those above it.

This topology is intentionally **one-directional**. Higher layers define meaning and boundaries; lower layers operate on those definitions without altering them. No layer substitutes for another, and no layer inherits the function of those above or below it. This structure ensures that representational clarity is maintained even as analytical or evaluative components evolve.

3. AI Perception — Disciplinary Layer

AI Perception constitutes the **disciplinary foundation** of the EntityWorks Standard. It defines the conceptual problem space concerned with how AI systems form, maintain, and express representations of people, organisations, relationships, and ideas.

This layer establishes the interpretive vocabulary and boundaries used throughout the Standard. It does not describe system internals, optimisation techniques, or behavioural prescriptions. Instead, it provides the **language through which representational behaviour can be described consistently**, forming the conceptual ground upon which all other components depend.

AI Interpretation & Reliance Domain — Operational Reliance Boundary

The **AI Interpretation & Reliance Domain** defines a distinct operational domain concerned with how AI-generated explanations, summaries, and descriptions are interpreted and relied upon by humans or institutions as shared context for understanding.

This domain exists **after representational content has been produced and expressed, but before that content is acted upon through decisions, policies, enforcement, or other consequential actions**. It concerns the point at which AI-expressed understanding is taken up, stabilised, and treated as sufficient basis for further reasoning.

The AI Interpretation & Reliance Domain does not describe how AI systems compute outputs, nor how decisions are executed. It does not evaluate representational quality, prescribe governance mechanisms, or assert authority over outcomes. Its scope is limited to the interpretive handoff at which AI-generated explanations become socially operative through reliance.

Within the EntityWorks body of work, this domain:

- Operates **within the scope of the discipline of AI Perception**, as it concerns the uptake of AI-expressed understanding.
- Remains **analytically distinct from the layered representational topology** of the EntityWorks Standard.
- Intersects with, but does not subsume, representational boundary artefacts such as Machine-Facing Pages.
- Provides contextual grounding for why diagnostic, evaluative, and analytical components become consequential once AI explanations are relied upon by others.

The AI Interpretation & Reliance Domain introduces no new mechanisms, criteria, or evaluative functions. It serves as a **structural boundary reference**, making explicit an operational domain that already exists in practice but has not previously been named in a stable or canonical form.

4. Entity Understanding Layer (EUL) — Representational Architecture

The Entity Understanding Layer articulates the **structural architecture** through which representational behaviour is described. It translates the abstract concerns of AI Perception into a coherent representational model, defining how entities, attributes, and relationships are organised within the scope of the Standard.

EUL does not evaluate or judge representational quality. It provides the architectural lens required for subsequent diagnostic and evaluative components to operate meaningfully. Without this architectural layer, later analysis would lack a stable reference point for interpretation.

5. Machine-Facing Pages (MFP) & Machine-Facing Page Declaration (MFPD)

Representational Boundary Layer

Machine-Facing Pages define a **class of digital surfaces** intended primarily for machine interpretation rather than human navigation. The Machine-Facing Page Declaration binds declared intent to those surfaces, making explicit the purpose and role of such pages within an organisation's publishing footprint.

This layer exists **exclusively at the boundary** between publication and machine interpretation. MFPs and the MFPD do not evaluate representational quality, enforce standards, or signal compliance. They define *where* representational interaction occurs and *what* that interaction is intended to be — nothing more.

Concrete example (abstracted):

An organisation may publish structured reference material for AI systems that is not linked from its human-facing site. The MFPD exists to state *why* that material exists and *what it is for*, without asserting whether it is good, correct, compliant, or aligned.

6. Failure Modes — Diagnostic Layer

Failure modes describe **recognisable conditions** under which representational behaviour departs from the structures defined by the Standard. These components provide diagnostic language for identifying instability, collision, or interpretive breakdowns without asserting causality at the system level.

Failure modes do not prescribe remediation or rank severity. They exist to support shared understanding when analysing representational behaviour, enabling consistent discussion of issues such as ambiguity, collapse, or drift across contexts.

7. Entity Discoverability Index (EDI) — Evaluative Layer

The Entity Discoverability Index operates at the **evaluative layer**, assessing aspects of representational clarity and distinguishability as defined within the Standard. It produces evaluative outputs based on predefined criteria, without analysing temporal behaviour or prescribing corrective action.

The EDI depends on upstream layers for meaning and scope. It does not alter architectural definitions, redefine boundaries, or interpret trajectories over time. Its role is strictly to evaluate representational conditions at a given point within the Standard's framework.

8. EntityWorks Analytics (EWA) — Analytical Layer

EntityWorks Analytics examines **representational conditions over time**, analysing how representations persist, change, or degrade across contexts and interactions. This layer focuses on patterns and trajectories rather than thresholds or pass/fail outcomes.

Analytics does not generate criteria, enforce standards, or issue signals. It relies on evaluative outputs and diagnostic structures defined upstream, transforming them into longitudinal insight without collapsing into governance or prescription.

9. AI Perception Integrity Mark (AIPM) — Conformance Signalling Layer

The AI Perception Integrity Mark functions as a **conformance signal** within the scope of the EntityWorks Standard. It indicates whether published representations meet the structural and interpretive criteria defined internally by the Standard.

The AIPM does not certify systems, enforce compliance, or govern behaviour. It operates downstream of all analytical and evaluative components, signalling alignment without asserting authority beyond the Standard's defined scope.

10. Non-Aggregability Rule

Components of the EntityWorks Standard are **not aggregatable across layers**. No component may perform the function of another, and no artefact may legitimately combine multiple layer roles into a single mechanism.

Any construct that attempts to merge boundary definition with evaluation, evaluation with governance, or analytics with signalling introduces **representational ambiguity** and falls outside the scope of the Standard. This rule preserves interpretive clarity by preventing functional collapse, even where such collapse might appear efficient or commercially attractive.

11. Canonical Structural Order

In structural order, the topology is as follows:

1. AI Perception — disciplinary foundation
2. Entity Understanding Layer — representational architecture
3. MFP / MFPD — representational boundary
4. Failure Modes — diagnostics
5. Entity Discoverability Index — evaluation
6. EntityWorks Analytics — analysis
7. AI Perception Integrity Mark — conformance signalling

Each layer depends on those above it and constrains those below it. No layer reverses this dependency.

12. Scope Limitation

This reference governs **structural interpretation only**. It does not assert authority over external systems, practices, governance frameworks, or regulatory regimes. Its purpose is to define how components of the EntityWorks Standard relate to one another, not how they are applied or enforced in external contexts.