

AI-Mediated Representation Risk (AMRR)

Canonical Definition

Published by EntityWorks

Document Control

Document Title: AI-Mediated Representation Risk (AMRR) — Canonical Definition

Edition: Definition

Version: v1.0

Status: Published

Part of: The EntityWorks Standard

Publication Date: January 2026

Maintained by: EntityWorks Ltd

Scope and Status Notice

This document constitutes an official publication of the EntityWorks Standard.

It defines **AI-Mediated Representation Risk (AMRR)** as a descriptive term used to name an existing condition of the AI-mediated informational environment.

This document is **descriptive and non-normative**. It does not prescribe behaviour, mandate controls, recommend mitigation, or assert regulatory authority. It defines a class of risk exposure only, not responsibility, fault, compliance, or remedy.

The scope of this document is limited to naming and bounding the condition under which AI systems generate, stabilise, or propagate representations of organisations that are relied upon by third parties.

Copyright and Use Notice

© 2026 EntityWorks Ltd. All rights reserved.

This document may be cited, referenced, and redistributed in unmodified form for informational, academic, regulatory, or evaluative purposes, provided attribution to EntityWorks is preserved.

No part of this document may be modified, reissued, or presented as an authoritative definition of AI-Mediated Representation Risk without explicit reference to its origin within the EntityWorks Standard.

Definition

AI-Mediated Representation Risk (AMRR) refers to the exposure that arises when artificial intelligence systems generate, stabilise, or propagate representations of an organisation that are treated as authoritative or decision-relevant by third parties, regardless of whether those representations were intended, endorsed, or produced by the organisation itself.

This risk exists independently of an organisation's direct use of AI systems and arises from the broader AI-mediated informational environment in which all organisations are represented and interpreted through automated synthesis, shaping perception, reliance, and decision-making.

Scope

AI-Mediated Representation Risk applies where:

- AI systems generate or infer representations of an organisation based on available digital signals, surfaces, or contextual data
- Those representations are relied upon by external actors for evaluation, comparison, attribution, or decision-making
- The resulting exposure exists irrespective of the organisation's awareness of, intent toward, or control over the generating systems

AMRR concerns the **existence and emergence** of representational exposure, not its management, mitigation, or prevention.

Exclusions

AI-Mediated Representation Risk does **not**:

- Describe AI safety, alignment, or ethical considerations
- Specify technical failures, model internals, or system performance characteristics
- Assign legal liability, fault, or responsibility
- Recommend governance actions, controls, remediation, or optimisation
- Provide advice, operational guidance, or implementation strategies

AMRR is a **descriptive term**, not a prescriptive framework.

EntityWorks' Role

EntityWorks defines AI-Mediated Representation Risk as part of its role as a **standards-level interpretive infrastructure organisation** concerned with how AI systems form, stabilise, and propagate representations of organisations.

EntityWorks does not provide mitigation, optimisation, or remediation services in relation to this risk.