

EntityWorks Component Relationship Map — Structural Reference (v2.0)

Edition: Human-Facing Structural Reference

Version: v2.0

Status: Published

Original Publication: 27 December 2025

Publication Date: 11 January 2026

Maintained by: EntityWorks Ltd

Abstract

This document provides a structured description of how the published components of the EntityWorks body of work relate to one another. It defines the **positional and dependency relationships** between disciplines, architectural layers, boundary artefacts, diagnostic constructs, evaluative mechanisms, analytical functions, and conformance signals as they exist across current EntityWorks publications.

This document is **descriptive only**. It introduces no new components, criteria, functions, or mechanisms.

Scope and Status Notice

This reference applies exclusively to components **explicitly defined and published by EntityWorks** at the time of release. It describes relationships between existing components as documented in their respective publications.

This document does **not**:

- define new standards or requirements
- prescribe implementation approaches
- combine or aggregate component functions
- describe technical system behaviour or internal model mechanisms

Its authority is limited to **structural interpretation** within the EntityWorks Standard.

Copyright and Use Notice

© 2025 EntityWorks Ltd. All rights reserved.

This document may be cited, referenced, and redistributed in unmodified form for informational, academic, regulatory, or evaluative purposes, provided attribution to EntityWorks Ltd is preserved.

No part of this document may be modified, reissued, or presented as an authoritative standard without explicit reference to its origin within the EntityWorks Standard.

1. Purpose of This Reference

This document establishes the **canonical relational topology** of the EntityWorks Standard.

Its purpose is to preserve **structural clarity** by defining how components relate to one another and, critically, how they **do not**. By explicitly separating concerns across disciplinary, architectural, boundary, diagnostic, evaluative, analytical, and signalling layers, this reference prevents functional overlap, role confusion, and scope inflation as the Standard is applied, discussed, or extended.

This document does not introduce mechanisms or criteria. It exists to ensure that each component is interpreted **only** according to its defined role and position within the overall structure.

2. Structural Principle: Layered Representational Topology

The EntityWorks Standard is organised as a **layered representational topology**, not as a unified framework or composite system.

Each layer addresses a distinct aspect of how entities are represented, interpreted, evaluated, analysed, or signalled. Layers are conceptually dependent on those above them and operate within boundaries defined upstream.

This topology is **strictly one-directional**:

- Higher layers define meaning, scope, and representational boundaries
- Lower layers operate on those definitions without altering them

No layer substitutes for another. No layer inherits the function of those above or below it. This structure ensures that representational clarity is preserved even as analytical or evaluative activity evolves.

3. AI Perception — Disciplinary Layer

AI Perception constitutes the **disciplinary foundation** of the EntityWorks Standard.

It defines the conceptual problem space concerned with how AI systems form, maintain, stabilise, and express representations of people, organisations, relationships, and ideas. This layer establishes the interpretive vocabulary and conceptual boundaries used throughout the Standard.

AI Perception does **not** describe system internals, optimisation techniques, behavioural prescriptions, or governance requirements. Its role is exclusively to provide a stable language for describing representational behaviour, forming the conceptual ground upon which all other components depend.

4. AI Interpretation & Reliance Domain — Operational Reliance Boundary

The AI Interpretation & Reliance Domain defines a distinct operational domain concerned with how AI-generated explanations, summaries, and descriptions are **interpreted and relied upon** by humans or institutions as shared context for understanding.

This domain exists **after representational content has been expressed** and **before consequential action** is taken. It concerns the point at which AI-expressed understanding is treated as sufficient basis for further reasoning, decision-making, or attribution.

This domain does not describe how AI systems compute outputs, evaluate representational quality, or execute decisions. It introduces no criteria, mechanisms, or authority. Its role is to **name and bound** an operational domain that already exists in practice but is often left implicit.

5. Entity Understanding Layer (EUL) — Representational Architecture

The Entity Understanding Layer defines the **representational architecture** through which AI-formed understanding is described within the Standard.

It translates the abstract concerns of AI Perception into a coherent architectural model, specifying how entities, attributes, and relationships are organised and stabilised across representational contexts.

EUL does not judge quality, assess correctness, or produce evaluative output. Its function is to provide the **structural reference frame** required for diagnostic, evaluative, and analytical components to operate meaningfully and consistently.

6. Machine-Facing Pages (MFP) & Machine-Facing Page Declaration (MFPD) — Representational Boundary Layer

Machine-Facing Pages define a class of digital surfaces interpreted primarily by AI systems rather than human audiences.

The Machine-Facing Page Declaration binds declared intent to those surfaces, making explicit their purpose and role within an organisation's publishing footprint.

This layer exists exclusively at the boundary between publication and machine interpretation. It does not evaluate representational quality, enforce standards, or signal alignment. Its role is to define **where representational interaction occurs and what it is intended to be**, without asserting value, correctness, or compliance.

7. Failure Modes — Diagnostic Layer

Failure modes describe recognisable conditions under which representational behaviour departs from the structures defined by the Standard.

They provide **diagnostic language**, not explanation or remediation. Failure modes enable shared understanding when analysing instability, collision, fragmentation, or interpretive breakdown, without asserting causality, ranking severity, or prescribing corrective action.

This layer exists to support consistent diagnosis while remaining analytically neutral.

8. Entity Discoverability Index (EDI) — Evaluative Layer

The Entity Discoverability Index operates at the **evaluative layer**, assessing aspects of representational clarity, separability, and interpretive consistency as defined by upstream components.

EDI produces evaluative outputs based on predefined criteria. It does not analyse temporal behaviour, infer trajectories, or propose remediation. Its role is strictly to evaluate representational conditions **at a given point** within the Standard's framework.

EDI depends entirely on upstream layers for meaning and scope and does not alter architectural definitions or boundary conditions.

9. EntityWorks Analytics (EWA) — Analytical Layer

EntityWorks Analytics operates at the **analytical layer**, examining representational behaviour **over time**.

It analyses persistence, change, drift, or degradation across contexts and interactions, transforming evaluative and diagnostic inputs into longitudinal insight.

This layer does not generate criteria, issue judgements, or signal alignment. It exists to identify patterns and trajectories while remaining distinct from governance, enforcement, or prescription.

10. AI Perception Integrity Mark (AIPM) — Conformance Signalling Layer

The AI Perception Integrity Mark functions as a **conformance signal** within the scope of the EntityWorks Standard.

It operates downstream of diagnostic, evaluative, and analytical components, signalling whether published representations meet defined structural and interpretive conditions.

The AIPM does not certify systems, enforce compliance, or govern behaviour. Its authority is limited to signalling alignment **within** the Standard's defined scope.

11. Non-Aggregatability Rule

Components of the EntityWorks Standard are **not aggregatable across layers**.

No component may perform the function of another, and no artefact may legitimately combine multiple layer roles into a single mechanism. Any construct that merges boundary definition, evaluation, analysis, or signalling introduces representational ambiguity and falls outside the scope of the Standard.

This rule preserves interpretive clarity by preventing functional collapse, even where such collapse may appear efficient or commercially attractive.

12. Canonical Structural Order

The canonical structural order of the EntityWorks Standard is:

1. AI Perception — disciplinary foundation
2. Entity Understanding Layer — representational architecture
3. MFP / MFPD — representational boundary
4. Failure Modes — diagnostics
5. Entity Discoverability Index — evaluation
6. EntityWorks Analytics — analysis
7. AI Perception Integrity Mark — conformance signalling

Each layer depends on those above it and constrains those below it. No layer reverses this dependency.

13. Scope Limitation

This reference governs **structural interpretation only**. It does not assert authority over external systems, organisational practices, governance frameworks, or regulatory regimes. Its purpose is to define how components of the EntityWorks Standard relate to one another, not how they are applied or enforced externally.