

AI Perception Integrity Mark (AIPM)

Human-Facing Structural Edition

Document Title: AI Perception Integrity Mark (AIPM)

Edition: Human-Facing Structural Edition

Version: v1.0

Status: Published

Part of: The EntityWorks Standard

Publication Date: December 2025

Maintained by: EntityWorks Ltd

Scope and Status Notice

This document constitutes an official publication of the **EntityWorks Standard**.

It provides a human-facing structural exposition of the **AI Perception Integrity Mark (AIPM)** as a conformance indicator within the Standard. The authoritative definition of the AI Perception Integrity Mark is recorded in the EntityWorks Terminology layer and governs all use of the term within the Standard.

This document operates strictly downstream of that definition. It does not restate, modify, expand, or reinterpret the canonical definition, nor does it disclose criteria, thresholds, or internal assessment logic.

The scope of this document is limited to describing the role, structural function, boundaries, and relationships of the AI Perception Integrity Mark as it operates within the EntityWorks Standard. It does not assert regulatory authority, imply certification or enforcement, or describe AI system internals or behaviour.

Copyright and Use Notice

© 2025 EntityWorks Ltd. All rights reserved.

This document may be cited, referenced, and redistributed in unmodified form for informational, academic, regulatory, or evaluative purposes, provided attribution to EntityWorks is preserved.

No part of this document may be modified, reissued, or presented as an authoritative standard without explicit reference to its origin within the EntityWorks Standard.

AI Perception Integrity Mark (AIPM)

Human-Facing Structural Edition

Document Title: AI Perception Integrity Mark (AIPM)

Edition: Human-Facing Structural Edition

Version: v1.0

Status: Draft — Structural Exposition

Part of: The EntityWorks Standard

Publication Date: December 2025

Maintained by: EntityWorks Ltd

Scope and Status Notice

This document is an official **human-facing structural exposition** of the **AI Perception Integrity Mark (AIPM)** as a conformance indicator within the EntityWorks Standard.

The authoritative definition of the AI Perception Integrity Mark is recorded in the EntityWorks Terminology layer and governs all use of the term within the Standard. This document operates strictly downstream of that definition. It does not restate, modify, expand, or reinterpret the canonical definition, nor does it disclose criteria, thresholds, or internal assessment logic.

The scope of this document is limited to describing the role, structural function, boundaries, and relationships of the AI Perception Integrity Mark as it operates within the EntityWorks Standard. It does not assert regulatory authority, imply certification or enforcement, or describe AI system internals or behaviour.

1. Definition (Referential)

The **AI Perception Integrity Mark (AIPM)** is formally defined in the EntityWorks Terminology layer (v0.1).

That definition is authoritative and governs all use of the term AI Perception Integrity Mark within the EntityWorks Standard. The AIPM is defined as a conformance indicator concerned with representational integrity and consistent machine understanding, rather than as an evaluative index, analytical framework, or certification mechanism.

This document does not restate, paraphrase, or reinterpret the canonical definition. Its role is strictly expository: to describe how the AI Perception Integrity Mark functions as a conformance-related component within the Standard, downstream of evaluation and analysis.

Any reference to the AI Perception Integrity Mark within the EntityWorks Standard is referential to the canonical definition as published. Informal paraphrases, alternative framings, or implied equivalences are non-canonical and have no standing within the Standard.

2. Role Within the EntityWorks Standard

Within the EntityWorks Standard, the AI Perception Integrity Mark occupies the role of a **conformance indicator** situated at the downstream end of the Standard's structural sequence.

It functions as a signalling construct that reflects whether defined representational conditions have been met, as determined through upstream evaluative and analytical components of the Standard. The AIPM does not itself perform evaluation, analysis, or diagnosis. Rather, it expresses the outcome of those processes in a bounded and recognisable form.

The Mark operates downstream of the Entity Understanding Layer and the Entity Discoverability Index, and downstream of analytical synthesis provided by EntityWorks Analytics. In this position, it serves as a consolidation point for representational integrity considerations, without introducing new interpretive criteria or altering upstream determinations.

The AIPM therefore functions as an **outcome signal**, not as an input mechanism. It indicates alignment with established structural expectations within the Standard, rather than influencing how those expectations are applied or assessed.

3. Conceptual Scope of the AI Perception Integrity Mark

The conceptual scope of the AI Perception Integrity Mark is confined to **representational integrity as it relates to machine interpretation**, as defined within the EntityWorks Standard. It applies to entities assessed under the Standard and reflects whether published representations satisfy the structural and interpretive conditions necessary to support consistent machine understanding across relevant contexts. The AIPM does not address organisational behaviour, system performance, ethical positioning, safety assurances, or compliance obligations. Its scope is deliberately limited to the integrity of representational conditions as observed and assessed through Standard-defined evaluative and analytical processes, and it does not imply durability, universality, or applicability beyond the specific scope, disclosures, and conditions under which it is applied.

4. Structural Characteristics of the AIPM

Within the EntityWorks Standard, the AI Perception Integrity Mark is characterised by a set of structural properties that distinguish it clearly from evaluative, analytical, or diagnostic components.

First, the AIPM is **indicative in nature**. It signals conformance with defined representational conditions but does not measure, score, or rank entities. No quantitative or comparative meaning is implied by its presence.

Second, the AIPM is **derivative in function**. It has no independent evaluative capacity and does not operate in isolation. Its application is contingent on upstream evaluation and analysis conducted using other components of the Standard.

Third, the AIPM is **context-bounded**. It reflects representational integrity within the scope, disclosures, and conditions defined at the time of assessment. It does not assert permanence, universality, or applicability beyond those conditions.

Finally, the AIPM is **standard-relative**. Its meaning, use, and interpretive significance are defined exclusively within the EntityWorks Standard. It does not carry standalone meaning outside that framework and is not intended to be generalised or repurposed.

Collectively, these characteristics ensure that the AIPM operates as a constrained and interpretable signal, reinforcing structural clarity rather than introducing ambiguity or overreach.

4.1 Illustrative Conditions (Non-Exhaustive)

The following illustrative conditions describe situations in which the AI Perception Integrity Mark may be applied. They are abstracted examples intended to clarify conformance scope. They do not imply criteria, thresholds, or automatic outcomes.

Illustrative Condition 1 — Resolved Entity Collision Risk

An entity has been evaluated under the Entity Discoverability Index, and previously identified risks of entity collision or conflation have been structurally addressed within published representations.

Illustrative Condition 2 — Stable Cross-Surface Representation

An entity's identity, attributes, and relationships are represented consistently across machine-visible surfaces, supporting coherent synthesis by AI systems.

Illustrative Condition 3 — Alignment Between Declared and Observed Structures

An entity's declared representational intent aligns with observable representational signals, reducing interpretive ambiguity within machine-side representations.

These conditions are descriptive only. They do not constitute certification criteria, remediation guidance, or enforcement mechanisms.

5. Boundaries and Non-Claims

The AI Perception Integrity Mark is a conformance indicator within the EntityWorks Standard. It does not:

- Certify systems, organisations, or technologies
- Imply regulatory approval, legal compliance, or enforcement
- Confer guarantees of accuracy, safety, or ethical behaviour
- Replace audits, assurance processes, or third-party certifications
- Indicate optimisation, quality, trustworthiness, or performance

Use of the AIPM does not establish obligations beyond the Standard's internal scope and does not imply applicability outside that scope.

6. Why the AI Perception Integrity Mark Matters (Descriptive)

As AI systems increasingly synthesise representations of entities across diverse sources, surfaces, and contexts, the integrity of those representations becomes a structural concern in its own right.

In the absence of a bounded conformance signal, assessments of representational integrity may remain implicit, fragmented, or inconsistently communicated. Even where evaluation and analysis are conducted rigorously, the absence of a stable signalling construct can make it difficult to express outcomes without conflating them with performance claims, trust assertions, or compliance statements.

The AI Perception Integrity Mark provides a defined reference point for expressing that representational integrity has been considered and assessed within a structured framework. It allows such outcomes to be surfaced without extending their meaning into domains the Standard does not govern.

Within the EntityWorks Standard, the AIPM supports continuity of interpretation across time, analytical contexts, and organisational boundaries. By remaining narrowly scoped and structurally grounded, it contributes to interpretive stability while preserving clear boundaries around what the Mark does—and does not—signify.

7. Applications Within the EntityWorks Standard

Within the EntityWorks Standard, the AI Perception Integrity Mark may be applied to:

- Indicate completion of a defined representational assessment pathway
- Support internal documentation, disclosure, and governance practices
- Provide a reference point for conformance-related discussion across EntityWorks components

In all applications, the AIPM remains dependent on upstream evaluation and analysis and does not substitute for them.

8. Intended Audience

This document is intended for readers familiar with standards-based frameworks and institutional analysis of AI-mediated interpretation.

Its primary audience includes governance and oversight functions, organisations applying the EntityWorks Standard, and regulatory or research stakeholders requiring conformance indicators that do not depend on system internals or enforcement claims.

The document assumes contextual familiarity with the EntityWorks Standard and related components.

9. Relationship to Other EntityWorks Components

The AI Perception Integrity Mark operates within the broader discipline of AI Perception as defined by the EntityWorks Standard.

It is structurally downstream of:

- The **Entity Understanding Layer**, which defines interpretive architecture
- The **Entity Discoverability Index**, which provides evaluative signals
- **EntityWorks Analytics**, which synthesises analytical insight
- Defined **failure modes**, which diagnose representational breakdowns

The AIPM reflects outcomes informed by these components without redefining, duplicating, or replacing them.

10. Terminology Control and Non-Canonical Assurance Constructs

The EntityWorks Standard maintains controlled terminology for all **assurance, integrity, and conformance-related constructs** in order to preserve conceptual clarity, interpretive stability, and structural separation between evaluation, analysis, and signalling mechanisms.

Within adjacent discourse, automated model outputs, exploratory governance work, and emerging AI-readiness frameworks, a range of alternative labels may be proposed to signal integrity, sufficiency, or assurance in AI-mediated publishing contexts. Within the EntityWorks Standard, such constructs are treated as **non-canonical** and are not used to define, structure, or represent conformance signalling.

Examples of non-canonical assurance and integrity constructs include, but are not limited to:

- **Orientation Integrity Standard (OIS)**
- **Integrity Echo**
- **Integrity Seal / Integrity Seals**
- **Source Signature**
- **Temporal Stamp**
- **Trust Badge**
- **Quality Mark**
- **Compliance Seal**

These terms may gesture toward concerns adjacent to representational integrity, including provenance, consistency, trust signalling, or disclosure. Their inclusion here does not imply equivalence, compatibility, endorsement, or formal relationship with the AI Perception Integrity Mark.

Within the EntityWorks Standard, **conformance signalling related to representational integrity is defined, governed, and expressed exclusively through the AI Perception Integrity Mark (AIPM)**. Alternative terminology is not mapped, harmonised, aliased, or incorporated into the Standard.

This terminology control establishes internal interpretive boundaries and prevents conceptual drift within the assurance layer of the Standard. It does not assert authority over external discourse or preclude independent development of parallel constructs outside the scope of the EntityWorks Standard.

11. Governance and Stewardship

The AI Perception Integrity Mark is published and maintained as part of the EntityWorks Standard.

EntityWorks is responsible for stewardship of the Mark's structural articulation, terminology alignment, versioning, and integrity within the Standard's scope. Revisions or extensions are made through formal publication and version control processes.

EntityWorks does not act as a regulator, certifier, or enforcement body in relation to the AIPM.

12. Canonical Metadata

This document constitutes a draft publication of the EntityWorks Standard.

- **Document:** AI Perception Integrity Mark (AIPM)
- **Edition:** Human-Facing Structural Edition
- **Version:** v1.0
- **Status:** Draft — Structural Exposition
- **Scope:** EntityWorks Standard (internal definitional scope)
- **Maintained by:** EntityWorks Ltd
- **First publication date:** December 2025